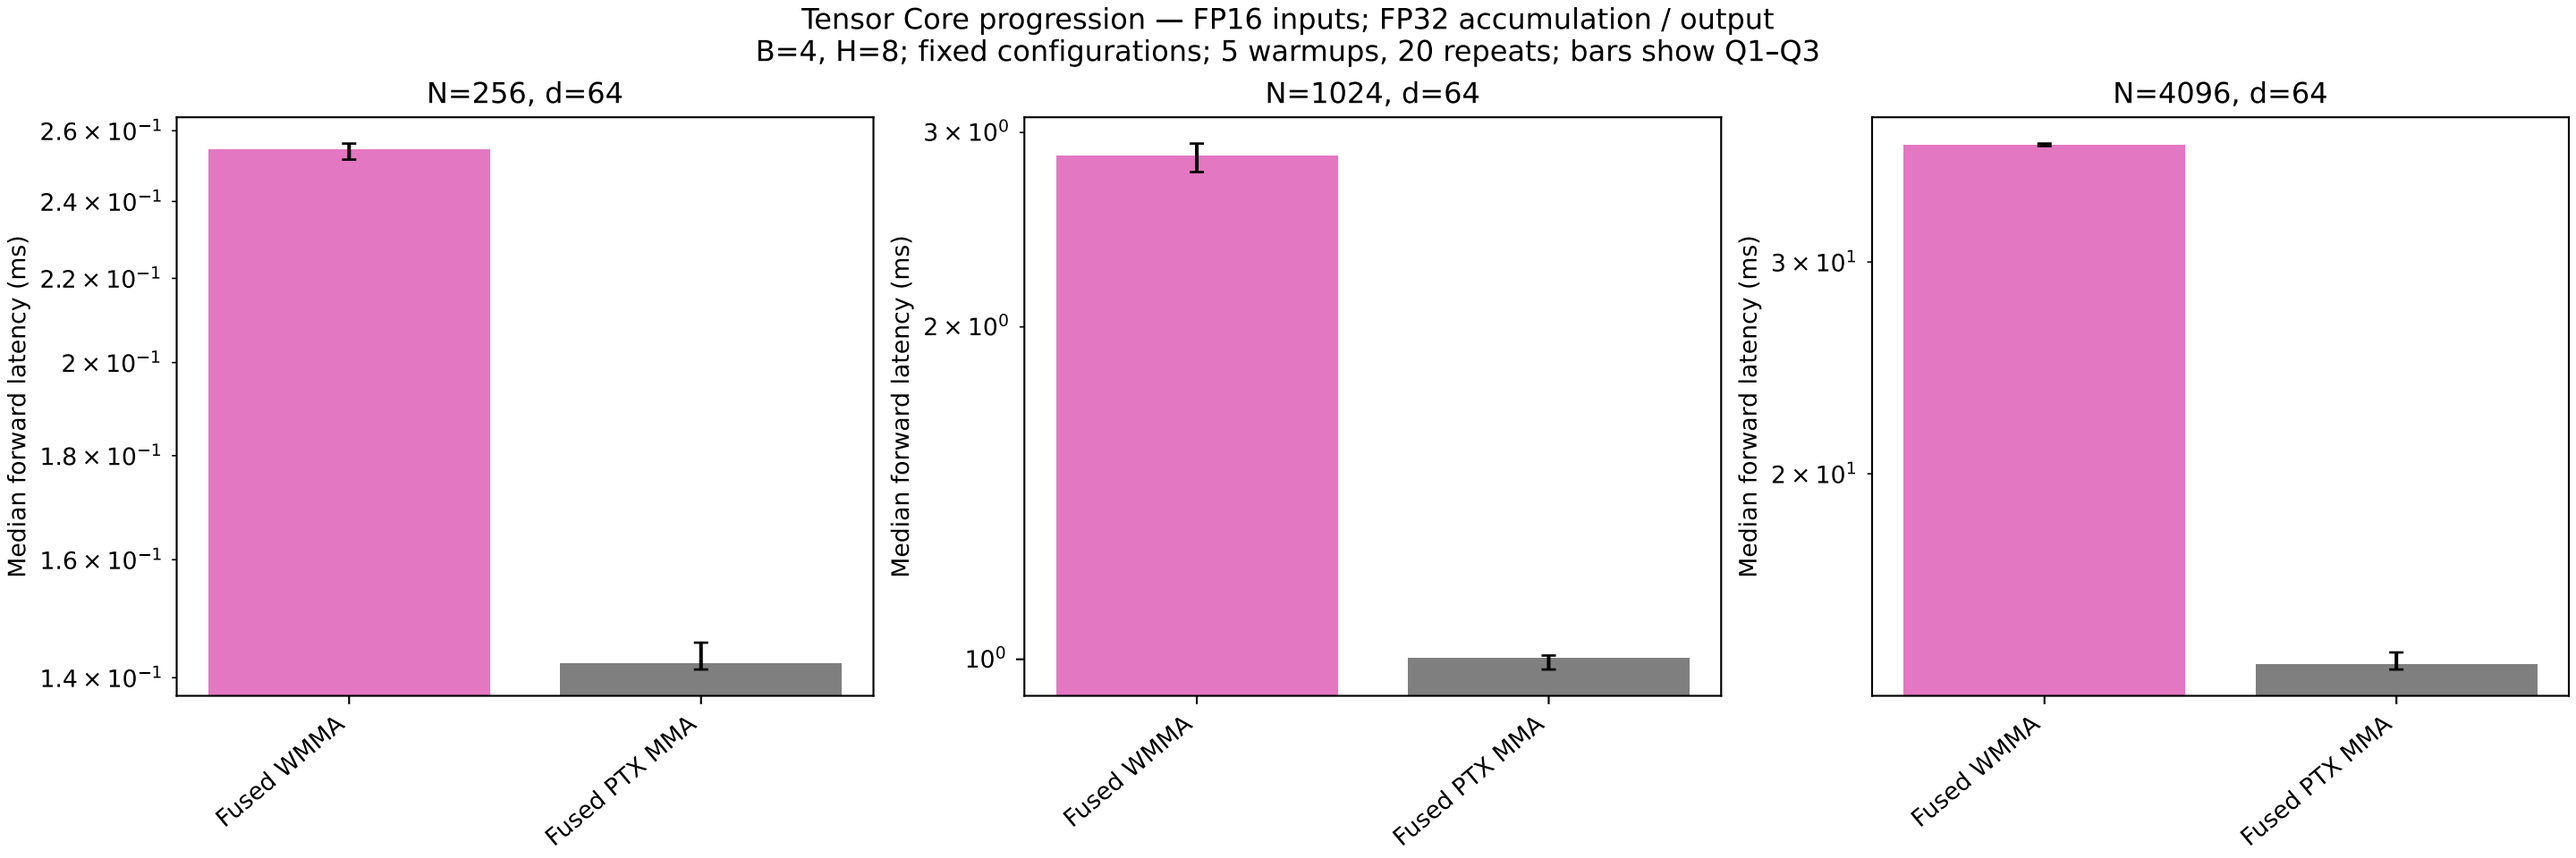
Tensor Core progression — FP16 inputs; FP32 accumulation / output
B=4, H=8; fixed configurations; 5 warmups, 20 repeats; bars show Q1–Q3
N=256, d=64
N=1024, d=64
N=4096, d=64
Median forward latency (ms)
2.6×10^{-1}
2.4×10^{-1}
2.2×10^{-1}
2×10^{-1}
1.8×10^{-1}
1.6×10^{-1}
1.4×10^{-1}
3×10^{0}
2×10^{0}
10^{0}
3×10^{1}
2×10^{1}
Fused WMMA
Fused PTX MMA