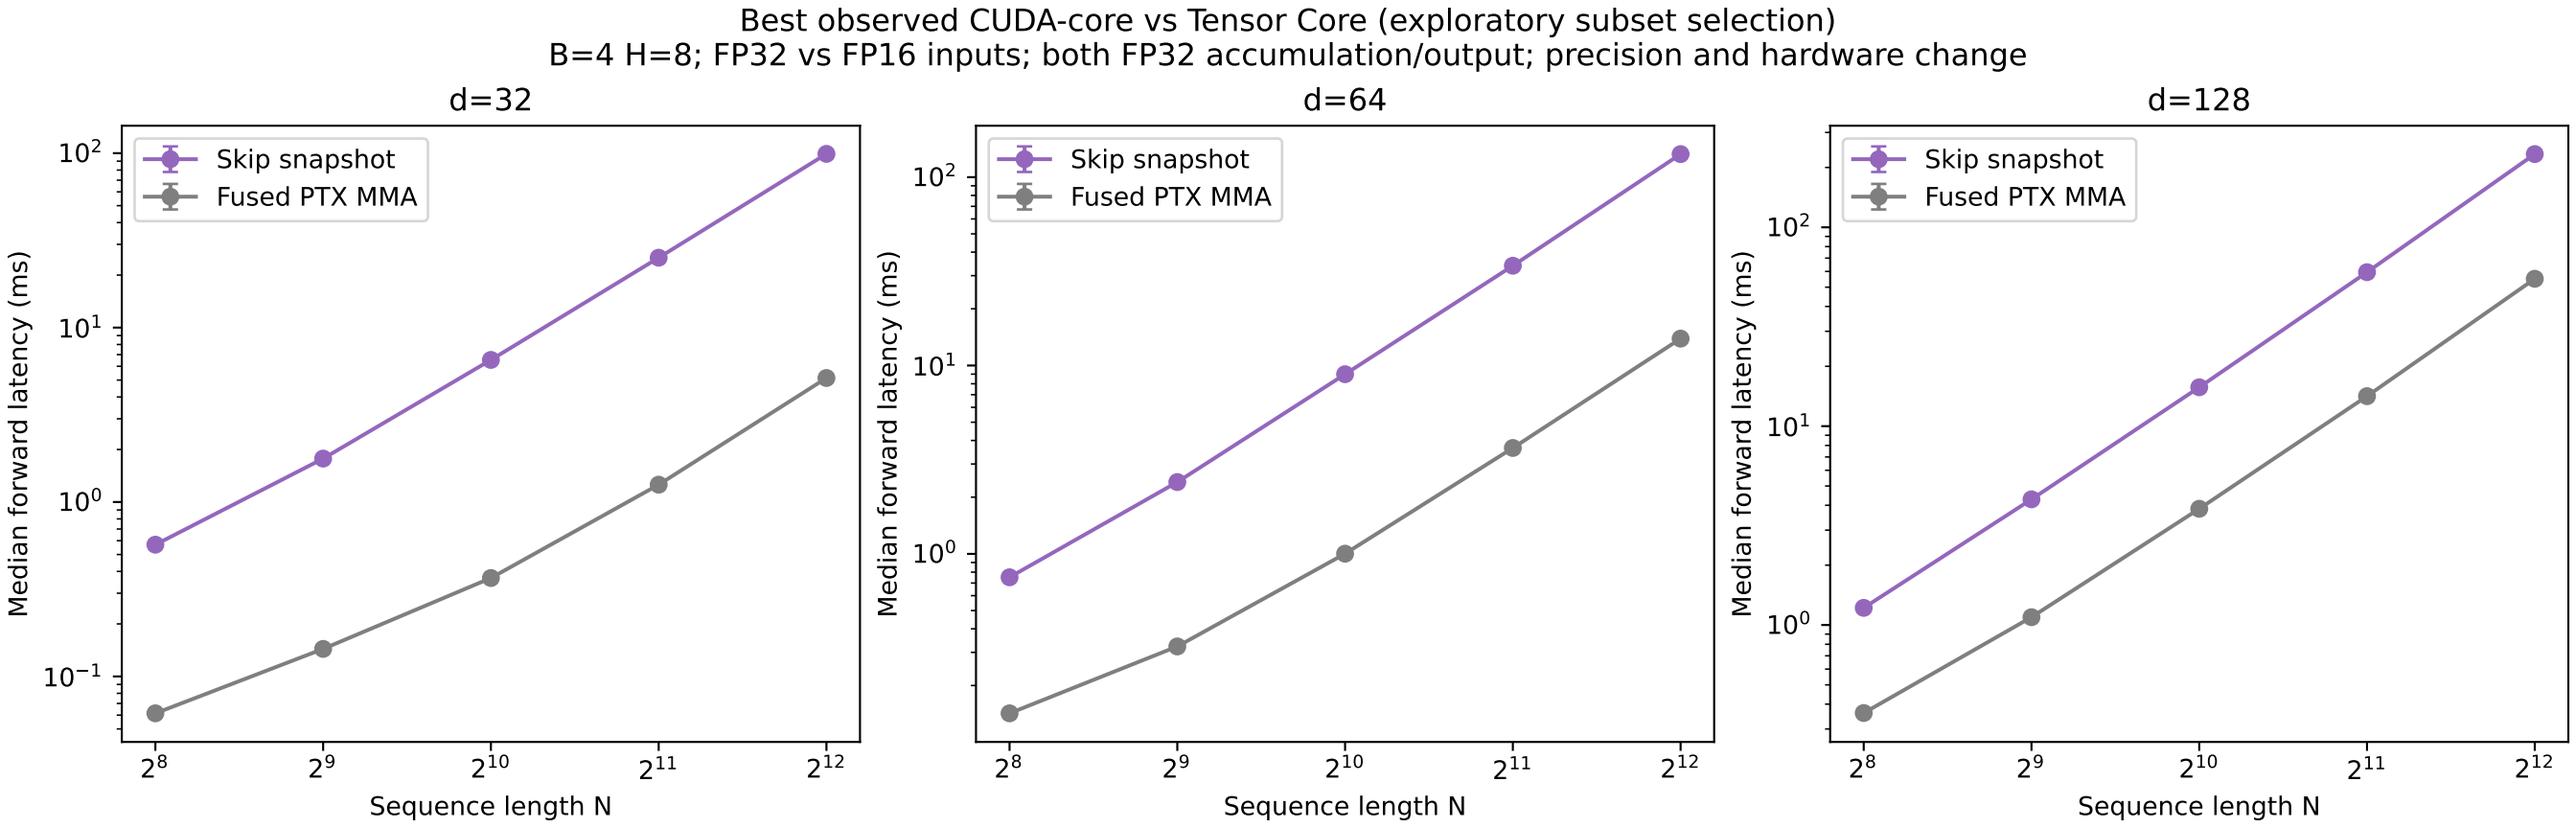
Best observed CUDA-core vs Tensor Core (exploratory subset selection)
B=4 H=8; FP32 vs FP16 inputs; both FP32 accumulation/output; precision and hardware change
d=32
d=64
d=128
Skip snapshot
Fused PTX MMA
Median forward latency (ms)
Sequence length N
$10^{2}$
$10^{1}$
$10^{0}$
$10^{-1}$
$2^{8}$
$2^{9}$
$2^{10}$
$2^{11}$
$2^{12}$